

Unsupervised Learning: Clustering

Finding Structure in Unlabeled Data

Nipun Batra

IIT Gandhinagar

October 29, 2025

What We'll Learn Today

1. **Why** unsupervised learning? (Motivation)
2. **What** is clustering? (Intuition first!)
3. **How** does K-Means work? (Step-by-step)
4. **When** does it fail? (Limitations)
5. **What** are alternatives? (Hierarchical, DBSCAN)

What We'll Learn Today

1. **Why** unsupervised learning? (Motivation)
2. **What** is clustering? (Intuition first!)
3. **How** does K-Means work? (Step-by-step)
4. **When** does it fail? (Limitations)
5. **What** are alternatives? (Hierarchical, DBSCAN)

Key Points: Key Points

Key Philosophy: We'll build intuition with examples, then add rigor!

Motivation

Supervised vs Unsupervised Learning

Supervised Learning

- Have: Features X + Labels Y
- Goal: Learn $f : X \rightarrow Y$
- Example: Spam detection
 - Email text \rightarrow Spam/Not Spam

Supervised vs Unsupervised Learning

Supervised Learning

- Have: Features X + Labels Y
- Goal: Learn $f : X \rightarrow Y$
- Example: Spam detection
 - Email text → Spam/Not Spam

Unsupervised Learning

- Have: Features X only (no labels!)
- Goal: Find **structure/patterns**
- Example: Customer segmentation
 - Customer data → Find groups

Supervised vs Unsupervised Learning

Supervised Learning

- Have: Features X + Labels Y
- Goal: Learn $f : X \rightarrow Y$
- Example: Spam detection
 - Email text → Spam/Not Spam

Unsupervised Learning

- Have: Features X only (no labels!)
- Goal: Find **structure/patterns**
- Example: Customer segmentation
 - Customer data → Find groups

Important:

Key Difference: We don't know the "right answer" beforehand!

Why Unsupervised Learning?

Three Main Reasons:

Why Unsupervised Learning?

Three Main Reasons:

1. Labels are expensive/impossible to get

- Medical images: Need expert radiologists (costly!)
- Customer behavior: No "true" groupings exist
- Exploratory analysis: Don't know what to look for yet

Why Unsupervised Learning?

Three Main Reasons:

1. Labels are expensive/impossible to get

- Medical images: Need expert radiologists (costly!)
- Customer behavior: No "true" groupings exist
- Exploratory analysis: Don't know what to look for yet

2. Discover hidden patterns

- Find new disease subtypes
- Identify market segments you didn't know existed
- Detect anomalies (fraud, network intrusion)

Why Unsupervised Learning?

Three Main Reasons:

1. Labels are expensive/impossible to get

- Medical images: Need expert radiologists (costly!)
- Customer behavior: No "true" groupings exist
- Exploratory analysis: Don't know what to look for yet

2. Discover hidden patterns

- Find new disease subtypes
- Identify market segments you didn't know existed
- Detect anomalies (fraud, network intrusion)

3. Data preprocessing

- Dimensionality reduction before supervised learning
- Feature extraction
- Data compression

Real-World Applications

Business

- **E-commerce:** Group customers by purchase behavior
- **Marketing:** Segment markets for targeted campaigns
- **Recommendation:**
"Customers who bought X also bought Y"

Real-World Applications

Business

- **E-commerce:** Group customers by purchase behavior
- **Marketing:** Segment markets for targeted campaigns
- **Recommendation:**
"Customers who bought X also bought Y"

Healthcare

- **Disease subtypes:** Find variants of cancer
- **Patient stratification:**
Personalized treatment

Real-World Applications

Business

- **E-commerce:** Group customers by purchase behavior
- **Marketing:** Segment markets for targeted campaigns
- **Recommendation:**
"Customers who bought X also bought Y"

Healthcare

- **Disease subtypes:** Find variants of cancer
- **Patient stratification:**
Personalized treatment

Technology

- **Image segmentation:** Divide image into regions
- **Document clustering:** Organize large text collections
- **Anomaly detection:** Find unusual patterns

Real-World Applications

Business

- **E-commerce:** Group customers by purchase behavior
- **Marketing:** Segment markets for targeted campaigns
- **Recommendation:**
"Customers who bought X also bought Y"

Healthcare

- **Disease subtypes:** Find variants of cancer
- **Patient stratification:**
Personalized treatment

Technology

- **Image segmentation:** Divide image into regions
- **Document clustering:** Organize large text collections
- **Anomaly detection:** Find unusual patterns

Science

- **Astronomy:** Classify galaxies
- **Climate:** Identify weather patterns
- **Biology:** Group similar

What is Clustering?

Clustering: The Intuition

Informal Definition:

Group similar objects together, separate dissimilar objects

Clustering: The Intuition

Informal Definition:

Group similar objects together, separate dissimilar objects

Key Questions:

1. What does "similar" mean? (Need a distance/similarity measure)

Clustering: The Intuition

Informal Definition:

Group similar objects together, separate dissimilar objects

Key Questions:

1. What does "similar" mean? (Need a distance/similarity measure)
2. How many groups? (May or may not know beforehand)

Clustering: The Intuition

Informal Definition:

Group similar objects together, separate dissimilar objects

Key Questions:

1. What does "similar" mean? (Need a distance/similarity measure)
2. How many groups? (May or may not know beforehand)
3. What shape are clusters? (Spherical? Elongated? Arbitrary?)

Clustering: The Intuition

Informal Definition:

Group similar objects together, separate dissimilar objects

Key Questions:

1. What does "similar" mean? (Need a distance/similarity measure)
2. How many groups? (May or may not know beforehand)
3. What shape are clusters? (Spherical? Elongated? Arbitrary?)

Example: Everyday Example

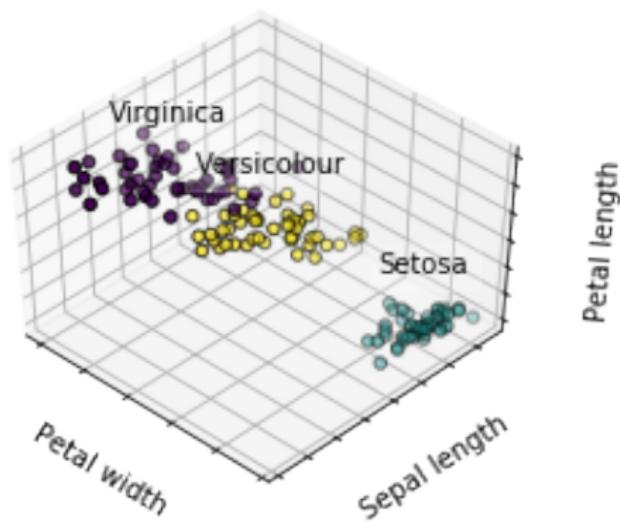
Organizing your wardrobe:

- **By color:** All red clothes together
- **By type:** All shirts together
- **By season:** All winter clothes together

Same clothes, different clusterings! Choice depends on your

Clustering: Visual Intuition

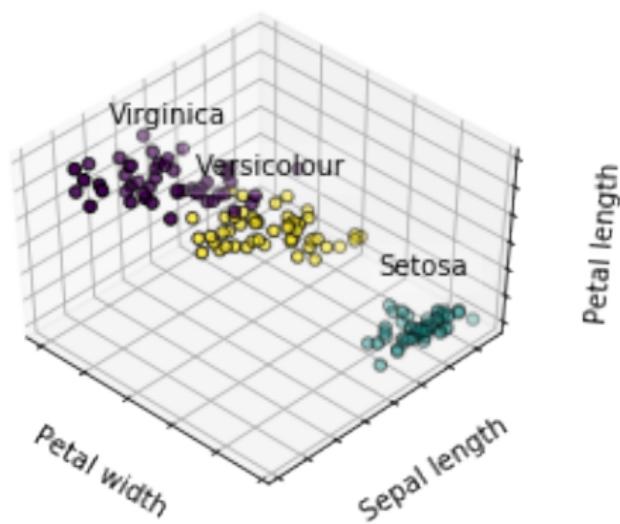
Ground Truth



Iris Dataset: Can you spot the groups?

Clustering: Visual Intuition

Ground Truth



Iris Dataset: Can you spot the groups?

- **Visual:** Easy to see 2-3 distinct groups
- **Question:** How do we make a computer see this?

Formal Problem Statement

Given:

- n data points: $\{x_1, x_2, \dots, x_n\}$ where $x_i \in \mathbb{R}^d$
- Number of clusters K (sometimes)
- Distance/similarity measure (usually Euclidean)

Formal Problem Statement

Given:

- n data points: $\{x_1, x_2, \dots, x_n\}$ where $x_i \in \mathbb{R}^d$
- Number of clusters K (sometimes)
- Distance/similarity measure (usually Euclidean)

Find:

- Partition data into K clusters: C_1, C_2, \dots, C_K
- Such that:
 - Points in same cluster are similar (small distances)
 - Points in different clusters are dissimilar (large distances)

Formal Problem Statement

Given:

- n data points: $\{x_1, x_2, \dots, x_n\}$ where $x_i \in \mathbb{R}^d$
- Number of clusters K (sometimes)
- Distance/similarity measure (usually Euclidean)

Find:

- Partition data into K clusters: C_1, C_2, \dots, C_K
- Such that:
 - Points in same cluster are similar (small distances)
 - Points in different clusters are dissimilar (large distances)

Mathematically:

- $C_1 \cup C_2 \cup \dots \cup C_K = \{1, 2, \dots, n\}$ (every point in some cluster)
- $C_i \cap C_j = \emptyset$ for $i \neq j$ (no overlaps)

Types of Clustering Algorithms

Definition (Clustering Algorithm Families)

Different approaches to finding clusters:

1. **Partitioning:** Divide data into K groups
 - K-Means, K-Medoids
 - Need to specify K

Types of Clustering Algorithms

Definition (Clustering Algorithm Families)

Different approaches to finding clusters:

1. **Partitioning:** Divide data into K groups
 - K-Means, K-Medoids
 - Need to specify K
2. **Hierarchical:** Build tree of clusters
 - Agglomerative (bottom-up), Divisive (top-down)
 - Don't need to specify K

Types of Clustering Algorithms

Definition (Clustering Algorithm Families)

Different approaches to finding clusters:

1. **Partitioning:** Divide data into K groups
 - K-Means, K-Medoids
 - Need to specify K
2. **Hierarchical:** Build tree of clusters
 - Agglomerative (bottom-up), Divisive (top-down)
 - Don't need to specify K
3. **Density-Based:** Find dense regions
 - DBSCAN, OPTICS
 - Can find arbitrary shapes

Types of Clustering Algorithms

Definition (Clustering Algorithm Families)

Different approaches to finding clusters:

1. **Partitioning:** Divide data into K groups
 - K-Means, K-Medoids
 - Need to specify K
2. **Hierarchical:** Build tree of clusters
 - Agglomerative (bottom-up), Divisive (top-down)
 - Don't need to specify K
3. **Density-Based:** Find dense regions
 - DBSCAN, OPTICS
 - Can find arbitrary shapes
4. **Model-Based:** Assume statistical model
 - Gaussian Mixture Models (GMM)
 - Probabilistic approach

Types of Clustering Algorithms

Definition (Clustering Algorithm Families)

Different approaches to finding clusters:

1. **Partitioning:** Divide data into K groups
 - K-Means, K-Medoids
 - Need to specify K
2. **Hierarchical:** Build tree of clusters
 - Agglomerative (bottom-up), Divisive (top-down)
 - Don't need to specify K
3. **Density-Based:** Find dense regions
 - DBSCAN, OPTICS
 - Can find arbitrary shapes
4. **Model-Based:** Assume statistical model
 - Gaussian Mixture Models (GMM)
 - Probabilistic approach

Today's Focus: K-Means (most popular), then Hierarchical and DBSCAN

K-Means: Intuition First!

K-Means: The Big Idea

Core Intuition:

*Each cluster has a **center point (centroid)**.*

Assign each point to the nearest centroid.

Update centroids based on assigned points.

Repeat until stable!

K-Means: The Big Idea

Core Intuition:

*Each cluster has a **center point (centroid)**.*

Assign each point to the nearest centroid.

Update centroids based on assigned points.

Repeat until stable!

Why "K-Means"?

- **K:** Number of clusters
- **Means:** Centroids are computed as means (averages)

K-Means: The Big Idea

Core Intuition:

*Each cluster has a **center point (centroid)**.*

Assign each point to the nearest centroid.

Update centroids based on assigned points.

Repeat until stable!

Why "K-Means"?

- **K:** Number of clusters
- **Means:** Centroids are computed as means (averages)

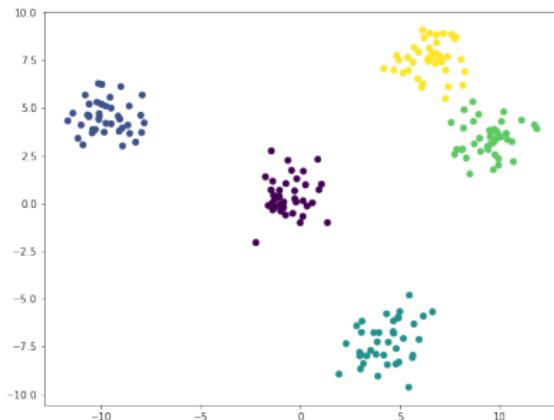
Key Points: Key Points

Two-Step Dance:

1. **Assignment:** Points → Nearest centroid
2. **Update:** Centroids → Mean of assigned points

Repeat until nothing changes!

K-Means: Visual Walkthrough (Step 0)



Raw data: Can you guess $K = 5$?

Start: Pick $K = 5$ centroids randomly (colored points)

K-Means: What Are We Optimizing?

Intuitive Goal:

*Make clusters **tight** - minimize distances within each cluster*

K-Means: What Are We Optimizing?

Intuitive Goal:

Make clusters **tight** - minimize distances within each cluster

Formal Objective: Minimize **Within-Cluster Sum of Squares (WCSS)**

$$\text{WCSS} = \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

where:

- $\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$ is the centroid of cluster k
- $\|x_i - \mu_k\|^2$ is squared distance from point to its centroid

K-Means: What Are We Optimizing?

Intuitive Goal:

Make clusters **tight** - minimize distances within each cluster

Formal Objective: Minimize **Within-Cluster Sum of Squares (WCSS)**

$$\text{WCSS} = \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

where:

- $\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$ is the centroid of cluster k
- $\|x_i - \mu_k\|^2$ is squared distance from point to its centroid

Important:

In words: Sum of squared distances from each point to its cluster center

Why This Objective Makes Sense

Minimizing WCSS means:

- Points close to their centroid → Tight clusters ✓

Why This Objective Makes Sense

Minimizing WCSS means:

- Points close to their centroid → Tight clusters ✓
- If WCSS is small → Clusters are compact ✓

Why This Objective Makes Sense

Minimizing WCSS means:

- Points close to their centroid → Tight clusters ✓
- If WCSS is small → Clusters are compact ✓
- Lower WCSS → Better clustering (usually) ✓

Why This Objective Makes Sense

Minimizing WCSS means:

- Points close to their centroid → Tight clusters ✓
- If WCSS is small → Clusters are compact ✓
- Lower WCSS → Better clustering (usually) ✓

Example: Concrete Example

Suppose we have 3 points in cluster C_1 : $(0, 0), (1, 0), (0, 1)$

- Centroid: $\mu_1 = (\frac{1}{3}, \frac{1}{3})$
- WCSS for C_1 :

$$\|(0, 0) - \mu_1\|^2 + \|(1, 0) - \mu_1\|^2 + \|(0, 1) - \mu_1\|^2 = \frac{2}{3}$$

K-Means Algorithm

K-Means Algorithm: The Recipe

Input:

- Data points $\{x_1, \dots, x_n\}$
- Number of clusters K

K-Means Algorithm: The Recipe

Input:

- Data points $\{x_1, \dots, x_n\}$
- Number of clusters K

Algorithm:

1. **Initialize:** Randomly choose K points as initial centroids $\{\mu_1, \dots, \mu_K\}$

K-Means Algorithm: The Recipe

Input:

- Data points $\{x_1, \dots, x_n\}$
- Number of clusters K

Algorithm:

1. **Initialize:** Randomly choose K points as initial centroids $\{\mu_1, \dots, \mu_K\}$
2. **Repeat** until convergence:

K-Means Algorithm: The Recipe

Input:

- Data points $\{x_1, \dots, x_n\}$
- Number of clusters K

Algorithm:

1. **Initialize:** Randomly choose K points as initial centroids $\{\mu_1, \dots, \mu_K\}$
2. **Repeat** until convergence:
 - 1) **Assignment Step:**
 - For each point x_i , assign to nearest centroid:

$$C_k^{(t)} = \{i : \|x_i - \mu_k^{(t)}\| \leq \|x_i - \mu_j^{(t)}\| \text{ for all } j\}$$

K-Means Algorithm: The Recipe

Input:

- Data points $\{x_1, \dots, x_n\}$
- Number of clusters K

Algorithm:

1. **Initialize:** Randomly choose K points as initial centroids $\{\mu_1, \dots, \mu_K\}$
2. **Repeat** until convergence:

1) Assignment Step:

- For each point x_i , assign to nearest centroid:

$$C_k^{(t)} = \{i : \|x_i - \mu_k^{(t)}\| \leq \|x_i - \mu_j^{(t)}\| \text{ for all } j\}$$

2) Update Step:

- Recompute each centroid as mean of assigned points:

$$\mu_k^{(t+1)} = \frac{1}{|C_k^{(t)}|} \sum_{i \in C_k^{(t)}} x_i$$

K-Means Algorithm: The Recipe

Input:

- Data points $\{x_1, \dots, x_n\}$
- Number of clusters K

Algorithm:

1. **Initialize:** Randomly choose K points as initial centroids $\{\mu_1, \dots, \mu_K\}$
2. **Repeat** until convergence:

1) Assignment Step:

- For each point x_i , assign to nearest centroid:

$$C_k^{(t)} = \{i : \|x_i - \mu_k^{(t)}\| \leq \|x_i - \mu_j^{(t)}\| \text{ for all } j\}$$

2) Update Step:

- Recompute each centroid as mean of assigned points:

$$\mu_k^{(t+1)} = \frac{1}{|C_k^{(t)}|} \sum_{i \in C_k^{(t)}} x_i$$

3. **Convergence:** Stop when assignments don't change (or change is tiny)

K-Means: Assignment Step in Detail

For each point x_i :

1. Compute distance to ALL K centroids:

$$d_k = \|x_i - \mu_k\|^2 \quad \text{for } k = 1, \dots, K$$

K-Means: Assignment Step in Detail

For each point x_i :

1. Compute distance to ALL K centroids:

$$d_k = \|x_i - \mu_k\|^2 \quad \text{for } k = 1, \dots, K$$

2. Find minimum distance:

$$k^* = \arg \min_k d_k$$

K-Means: Assignment Step in Detail

For each point x_i :

1. Compute distance to ALL K centroids:

$$d_k = \|x_i - \mu_k\|^2 \quad \text{for } k = 1, \dots, K$$

2. Find minimum distance:

$$k^* = \arg \min_k d_k$$

3. Assign x_i to cluster k^* :

$$\text{label}[i] = k^*$$

K-Means: Assignment Step in Detail

For each point x_i :

1. Compute distance to ALL K centroids:

$$d_k = \|x_i - \mu_k\|^2 \quad \text{for } k = 1, \dots, K$$

2. Find minimum distance:

$$k^* = \arg \min_k d_k$$

3. Assign x_i to cluster k^* :

$$\text{label}[i] = k^*$$

Example: Example

Point $x = (2, 3)$, centroids $\mu_1 = (1, 1)$, $\mu_2 = (4, 4)$

- $d_1 \equiv (2 - 1)^2 + (3 - 1)^2 \equiv 5$

K-Means: Update Step in Detail

For each cluster C_k :

1. Collect all points assigned to cluster k :

$$C_k = \{x_i : \text{label}[i] = k\}$$

K-Means: Update Step in Detail

For each cluster C_k :

1. Collect all points assigned to cluster k :

$$C_k = \{x_i : \text{label}[i] = k\}$$

2. Compute mean across all dimensions:

$$\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$$

K-Means: Update Step in Detail

For each cluster C_k :

1. Collect all points assigned to cluster k :

$$C_k = \{x_i : \text{label}[i] = k\}$$

2. Compute mean across all dimensions:

$$\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$$

3. This new μ_k becomes the centroid for next iteration

K-Means: Update Step in Detail

For each cluster C_k :

1. Collect all points assigned to cluster k :

$$C_k = \{x_i : \text{label}[i] = k\}$$

2. Compute mean across all dimensions:

$$\mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$$

3. This new μ_k becomes the centroid for next iteration

Example: Example

Cluster C_1 has points: $(0, 0), (2, 0), (0, 2)$

$$\mu_1 = \frac{1}{3} [(0, 0) + (2, 0) + (0, 2)] = \left(\frac{2}{3}, \frac{2}{3}\right)$$

K-Means: Worked Example (1/6)

Dataset: 6 points in 2D

Point	Coordinates
x_1	(1, 1)
x_2	(2, 1)
x_3	(4, 3)
x_4	(5, 4)
x_5	(1, 3)
x_6	(2, 2)

Goal: Cluster into $K = 2$ groups

Step 0: Initialize centroids randomly

- $\mu_1 = (1, 1)$ (pick x_1)
- $\mu_2 = (5, 4)$ (pick x_4)

K-Means: Worked Example (2/6) - Iteration 1: Assignment

Current centroids: $\mu_1 = (1, 1)$, $\mu_2 = (5, 4)$

Assign each point to nearest centroid:

Point	$d(\mu_1)$	$d(\mu_2)$	Nearest	Cluster
$x_1 = (1, 1)$	0	$\sqrt{25} = 5$	μ_1	C_1
$x_2 = (2, 1)$	1	$\sqrt{18} \approx 4.2$	μ_1	C_1
$x_3 = (4, 3)$	$\sqrt{13} \approx 3.6$	$\sqrt{2} \approx 1.4$	μ_2	C_2
$x_4 = (5, 4)$	5	0	μ_2	C_2
$x_5 = (1, 3)$	2	$\sqrt{17} \approx 4.1$	μ_1	C_1
$x_6 = (2, 2)$	$\sqrt{2} \approx 1.4$	$\sqrt{13} \approx 3.6$	μ_1	C_1

Result: $C_1 = \{x_1, x_2, x_5, x_6\}$, $C_2 = \{x_3, x_4\}$

K-Means: Worked Example (3/6) - Iteration 1: Update

Update centroids as means:

Cluster 1: $C_1 = \{(1, 1), (2, 1), (1, 3), (2, 2)\}$

$$\mu_1^{\text{new}} = \frac{1}{4}[(1, 1) + (2, 1) + (1, 3) + (2, 2)] = \frac{1}{4}(6, 7) = (1.5, 1.75)$$

K-Means: Worked Example (3/6) - Iteration 1: Update

Update centroids as means:

Cluster 1: $C_1 = \{(1, 1), (2, 1), (1, 3), (2, 2)\}$

$$\mu_1^{\text{new}} = \frac{1}{4}[(1, 1) + (2, 1) + (1, 3) + (2, 2)] = \frac{1}{4}(6, 7) = (1.5, 1.75)$$

Cluster 2: $C_2 = \{(4, 3), (5, 4)\}$

$$\mu_2^{\text{new}} = \frac{1}{2}[(4, 3) + (5, 4)] = \frac{1}{2}(9, 7) = (4.5, 3.5)$$

K-Means: Worked Example (3/6) - Iteration 1: Update

Update centroids as means:

Cluster 1: $C_1 = \{(1, 1), (2, 1), (1, 3), (2, 2)\}$

$$\mu_1^{\text{new}} = \frac{1}{4}[(1, 1) + (2, 1) + (1, 3) + (2, 2)] = \frac{1}{4}(6, 7) = (1.5, 1.75)$$

Cluster 2: $C_2 = \{(4, 3), (5, 4)\}$

$$\mu_2^{\text{new}} = \frac{1}{2}[(4, 3) + (5, 4)] = \frac{1}{2}(9, 7) = (4.5, 3.5)$$

Important:

Notice: Centroids moved from $(1, 1), (5, 4)$ to $(1.5, 1.75), (4.5, 3.5)$

They shifted toward the "center of mass" of their clusters!

K-Means: Worked Example (4/6) - Iteration 2: Assignment

New centroids: $\mu_1 = (1.5, 1.75)$, $\mu_2 = (4.5, 3.5)$

Reassign points:

Point	$d(\mu_1)$	$d(\mu_2)$	Nearest	Cluster
$x_1 = (1, 1)$	0.9	6.8	μ_1	C_1
$x_2 = (2, 1)$	0.8	7.1	μ_1	C_1
$x_3 = (4, 3)$	3.8	0.4	μ_2	C_2
$x_4 = (5, 4)$	6.8	0.4	μ_2	C_2
$x_5 = (1, 3)$	1.6	3.9	μ_1	C_1
$x_6 = (2, 2)$	0.5	6.6	μ_1	C_1

Result: Same as before! $C_1 = \{x_1, x_2, x_5, x_6\}$, $C_2 = \{x_3, x_4\}$

Convergence: Assignments didn't change → STOP!

K-Means: Worked Example (5/6) - Final Clustering

Final clusters:

- **Cluster 1:** $(1, 1), (2, 1), (1, 3), (2, 2)$ with centroid $(1.5, 1.75)$
- **Cluster 2:** $(4, 3), (5, 4)$ with centroid $(4.5, 3.5)$

K-Means: Worked Example (5/6) - Final Clustering

Final clusters:

- **Cluster 1:** (1, 1), (2, 1), (1, 3), (2, 2) with centroid (1.5, 1.75)
- **Cluster 2:** (4, 3), (5, 4) with centroid (4.5, 3.5)

WCSS calculation:

$$\begin{aligned} \text{WCSS} &= \sum_{i \in C_1} \|x_i - \mu_1\|^2 + \sum_{i \in C_2} \|x_i - \mu_2\|^2 \\ &= [0.9 + 0.8 + 1.6 + 0.5] + [0.4 + 0.4] \\ &= 3.8 + 0.8 = 4.6 \end{aligned}$$

K-Means: Worked Example (5/6) - Final Clustering

Final clusters:

- **Cluster 1:** (1, 1), (2, 1), (1, 3), (2, 2) with centroid (1.5, 1.75)
- **Cluster 2:** (4, 3), (5, 4) with centroid (4.5, 3.5)

WCSS calculation:

$$\begin{aligned} \text{WCSS} &= \sum_{i \in C_1} \|x_i - \mu_1\|^2 + \sum_{i \in C_2} \|x_i - \mu_2\|^2 \\ &= [0.9 + 0.8 + 1.6 + 0.5] + [0.4 + 0.4] \\ &= 3.8 + 0.8 = 4.6 \end{aligned}$$

Key Points: Key Points

Converged in 2 iterations!

Typically converges quickly (5-10 iterations), but worst-case can be slow.

K-Means: Worked Example (6/6) - Visualization

Visual summary:

- **Initial:** Random centroids far from optimal
- **Iteration 1:** Centroids move toward cluster centers
- **Iteration 2:** Centroids stabilize, assignments don't change
- **Final:** Two clear, compact clusters

K-Means: Worked Example (6/6) - Visualization

Visual summary:

- **Initial:** Random centroids far from optimal
- **Iteration 1:** Centroids move toward cluster centers
- **Iteration 2:** Centroids stabilize, assignments don't change
- **Final:** Two clear, compact clusters

Example: Key Insight

K-Means is iteratively refining the clustering:

1. Assignment makes WCSS smaller (each point goes to nearest centroid)
2. Update makes WCSS smaller (centroid is optimal for its assigned points)
3. Repeat until can't improve further (local optimum)

K-Means: Mathematical Rigor

Why K-Means Converges: Intuitive Proof

Claim: K-Means always converges

Why K-Means Converges: Intuitive Proof

Claim: K-Means always converges

Intuition: Each step decreases (or keeps same) the objective function WCSS

Why K-Means Converges: Intuitive Proof

Claim: K-Means always converges

Intuition: Each step decreases (or keeps same) the objective function WCSS

Proof sketch:

1. Assignment step:

- Each point assigned to **nearest** centroid
- Cannot increase WCSS (optimal choice given fixed centroids)

Why K-Means Converges: Intuitive Proof

Claim: K-Means always converges

Intuition: Each step decreases (or keeps same) the objective function WCSS

Proof sketch:

1. Assignment step:

- Each point assigned to **nearest** centroid
- Cannot increase WCSS (optimal choice given fixed centroids)

2. Update step:

- Centroid = mean of assigned points
- Mean minimizes sum of squared distances (calculus!)
- Cannot increase WCSS (optimal centroid given fixed assignments)

Why K-Means Converges: Intuitive Proof

Claim: K-Means always converges

Intuition: Each step decreases (or keeps same) the objective function WCSS

Proof sketch:

1. Assignment step:

- Each point assigned to **nearest** centroid
- Cannot increase WCSS (optimal choice given fixed centroids)

2. Update step:

- Centroid = mean of assigned points
- Mean minimizes sum of squared distances (calculus!)
- Cannot increase WCSS (optimal centroid given fixed assignments)

3. Conclusion:

- WCSS decreases or stays same at each step
- $\text{WCSS} \geq 0$ (bounded below)
- Finite number of possible assignments (K^n)
- Therefore, must converge! ✓

Why Mean Minimizes Sum of Squared Distances

Claim: For points $\{x_1, \dots, x_m\}$, the mean $\bar{x} = \frac{1}{m} \sum_{i=1}^m x_i$ minimizes $\sum_{i=1}^m \|x_i - c\|^2$

Why Mean Minimizes Sum of Squared Distances

Claim: For points $\{x_1, \dots, x_m\}$, the mean $\bar{x} = \frac{1}{m} \sum_{i=1}^m x_i$ minimizes $\sum_{i=1}^m \|x_i - c\|^2$

Proof: Let $f(c) = \sum_{i=1}^m \|x_i - c\|^2$. Taking derivative and setting to zero:

$$\frac{\partial f}{\partial c} = \sum_{i=1}^m 2(x_i - c)(-1) = 0$$

$$\sum_{i=1}^m (c - x_i) = 0$$

$$mc = \sum_{i=1}^m x_i$$

$$c = \frac{1}{m} \sum_{i=1}^m x_i = \bar{x} \quad \checkmark$$

Why Mean Minimizes Sum of Squared Distances

Claim: For points $\{x_1, \dots, x_m\}$, the mean $\bar{x} = \frac{1}{m} \sum_{i=1}^m x_i$ minimizes $\sum_{i=1}^m \|x_i - c\|^2$

Proof: Let $f(c) = \sum_{i=1}^m \|x_i - c\|^2$. Taking derivative and setting to zero:

$$\frac{\partial f}{\partial c} = \sum_{i=1}^m 2(x_i - c)(-1) = 0$$

$$\sum_{i=1}^m (c - x_i) = 0$$

$$mc = \sum_{i=1}^m x_i$$

$$c = \frac{1}{m} \sum_{i=1}^m x_i = \bar{x} \quad \checkmark$$

Key Points: Key Points

This is why we use **means** in K-Means - they're optimal!!

K-Means Objective: Equivalent Forms

Form 1: Within-Cluster Sum of Squares

$$\min \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

K-Means Objective: Equivalent Forms

Form 1: Within-Cluster Sum of Squares

$$\min \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

Form 2: Pairwise distances within clusters

$$\min \sum_{k=1}^K \frac{1}{|C_k|} \sum_{x_i \in C_k} \sum_{x_j \in C_k} \|x_i - x_j\|^2$$

K-Means Objective: Equivalent Forms

Form 1: Within-Cluster Sum of Squares

$$\min \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

Form 2: Pairwise distances within clusters

$$\min \sum_{k=1}^K \frac{1}{|C_k|} \sum_{x_i \in C_k} \sum_{x_j \in C_k} \|x_i - x_j\|^2$$

Equivalence:

$$\sum_{x_i \in C_k} \sum_{x_j \in C_k} \|x_i - x_j\|^2 = 2|C_k| \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

K-Means Objective: Equivalent Forms

Form 1: Within-Cluster Sum of Squares

$$\min \sum_{k=1}^K \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

Form 2: Pairwise distances within clusters

$$\min \sum_{k=1}^K \frac{1}{|C_k|} \sum_{x_i \in C_k} \sum_{x_j \in C_k} \|x_i - x_j\|^2$$

Equivalence:

$$\sum_{x_i \in C_k} \sum_{x_j \in C_k} \|x_i - x_j\|^2 = 2|C_k| \sum_{x_i \in C_k} \|x_i - \mu_k\|^2$$

Definition (Within-Cluster Variation)

Both forms measure how "tight" clusters are - smaller is better!

K-Means: Complexity Analysis

Time Complexity per iteration:

- **Assignment:** $O(nKd)$
 - For each of n points
 - Compute distance to K centroids
 - Each distance is $O(d)$ (dimensionality)
- **Update:** $O(nd)$
 - Sum up n points across d dimensions
 - Divide by cluster sizes

K-Means: Complexity Analysis

Time Complexity per iteration:

- **Assignment:** $O(nKd)$
 - For each of n points
 - Compute distance to K centroids
 - Each distance is $O(d)$ (dimensionality)
- **Update:** $O(nd)$
 - Sum up n points across d dimensions
 - Divide by cluster sizes

Total: $O(nKdt)$ where t = number of iterations

K-Means: Complexity Analysis

Time Complexity per iteration:

- **Assignment:** $O(nKd)$
 - For each of n points
 - Compute distance to K centroids
 - Each distance is $O(d)$ (dimensionality)
- **Update:** $O(nd)$
 - Sum up n points across d dimensions
 - Divide by cluster sizes

Total: $O(nKdt)$ where t = number of iterations

Space Complexity: $O((n + K)d)$

- Store n data points and K centroids

K-Means: Complexity Analysis

Time Complexity per iteration:

- **Assignment:** $O(nKd)$
 - For each of n points
 - Compute distance to K centroids
 - Each distance is $O(d)$ (dimensionality)
- **Update:** $O(nd)$
 - Sum up n points across d dimensions
 - Divide by cluster sizes

Total: $O(nKdt)$ where t = number of iterations

Space Complexity: $O((n + K)d)$

- Store n data points and K centroids

Key Points: Key Points

Practical: Usually converges in < 10 iterations → Very fast!

Choosing K: The Elbow Method

How Many Clusters?

Problem: K-Means requires specifying K beforehand

How Many Clusters?

Problem: K-Means requires specifying K beforehand

Question: How do we choose K ?

How Many Clusters?

Problem: K-Means requires specifying K beforehand

Question: How do we choose K ?

Bad idea: Just minimize WCSS

- $K = 1$: WCSS is large (one big cluster)
- $K = n$: WCSS is zero (each point is its own cluster)
- WCSS always decreases as K increases!

How Many Clusters?

Problem: K-Means requires specifying K beforehand

Question: How do we choose K ?

Bad idea: Just minimize WCSS

- $K = 1$: WCSS is large (one big cluster)
- $K = n$: WCSS is zero (each point is its own cluster)
- WCSS always decreases as K increases!

Important:

Trade-off: We want low WCSS but also not too many clusters!

The Elbow Method: Intuition

Idea: Plot WCSS vs K and look for an "elbow"

The Elbow Method: Intuition

Idea: Plot WCSS vs K and look for an "elbow"

Elbow = point where:

- Adding more clusters doesn't help much
- Diminishing returns set in
- Curve flattens out

The Elbow Method: Intuition

Idea: Plot WCSS vs K and look for an "elbow"

Elbow = point where:

- Adding more clusters doesn't help much
- Diminishing returns set in
- Curve flattens out

Analogy: Hiring employees

- $1 \rightarrow 2$ employees: Huge productivity gain
- $2 \rightarrow 3$ employees: Still helpful
- $10 \rightarrow 11$ employees: Marginal benefit
- $100 \rightarrow 101$ employees: Almost no difference

The Elbow Method: Intuition

Idea: Plot WCSS vs K and look for an "elbow"

Elbow = point where:

- Adding more clusters doesn't help much
- Diminishing returns set in
- Curve flattens out

Analogy: Hiring employees

- $1 \rightarrow 2$ employees: Huge productivity gain
- $2 \rightarrow 3$ employees: Still helpful
- $10 \rightarrow 11$ employees: Marginal benefit
- $100 \rightarrow 101$ employees: Almost no difference

Key Points: Key Points

Elbow = sweet spot where benefit of adding clusters drops sharply

Elbow Method: Mathematical Formulation

Algorithm:

1. Run K-Means for $K = 1, 2, 3, \dots, K_{\max}$
2. For each K , compute WCSS(K)
3. Plot WCSS(K) vs K
4. Find "elbow" = maximum curvature point

Elbow Method: Mathematical Formulation

Algorithm:

1. Run K-Means for $K = 1, 2, 3, \dots, K_{\max}$
2. For each K , compute $\text{WCSS}(K)$
3. Plot $\text{WCSS}(K)$ vs K
4. Find "elbow" = maximum curvature point

Finding elbow programmatically:

- Compute second derivative: $\frac{d^2(\text{WCSS})}{dK^2}$
- Elbow = maximum of second derivative

Elbow Method: Mathematical Formulation

Algorithm:

1. Run K-Means for $K = 1, 2, 3, \dots, K_{\max}$
2. For each K , compute $\text{WCSS}(K)$
3. Plot $\text{WCSS}(K)$ vs K
4. Find "elbow" = maximum curvature point

Finding elbow programmatically:

- Compute second derivative: $\frac{d^2(\text{WCSS})}{dK^2}$
- Elbow = maximum of second derivative

Or use heuristic:

$$\text{Elbow} = \arg \max_K [\text{WCSS}(K - 1) - \text{WCSS}(K)]$$

(biggest drop in WCSS)

Elbow Method: Example

Example WCSS vs K:

K	WCSS
1	1000
2	600 (drop: 400)
3	400 (drop: 200)
4	300 (drop: 100) ← Elbow!
5	250 (drop: 50)
6	220 (drop: 30)

Elbow Method: Example

Example WCSS vs K:

K	WCSS
1	1000
2	600 (drop: 400)
3	400 (drop: 200)
4	300 (drop: 100) ← Elbow!
5	250 (drop: 50)
6	220 (drop: 30)

- Drops slow down after $K = 4$
- Suggests $K = 4$ is optimal

Elbow Method: Example

Example WCSS vs K:

K	WCSS
1	1000
2	600 (drop: 400)
3	400 (drop: 200)
4	300 (drop: 100) ← Elbow!
5	250 (drop: 50)
6	220 (drop: 30)

- Drops slow down after $K = 4$
- Suggests $K = 4$ is optimal

Important:

Note: Elbow method is a **heuristic**, not a theorem!
Use domain knowledge + elbow method together.

Other Methods to Choose K

Silhouette Score:

- Measures how similar a point is to its own cluster vs other clusters
- Range: $[-1, 1]$, higher is better
- Choose K with highest average silhouette score

Other Methods to Choose K

Silhouette Score:

- Measures how similar a point is to its own cluster vs other clusters
- Range: $[-1, 1]$, higher is better
- Choose K with highest average silhouette score

Gap Statistic:

- Compare WCSS to expected WCSS under null (random) data
- Choose K where gap is largest

Other Methods to Choose K

Silhouette Score:

- Measures how similar a point is to its own cluster vs other clusters
- Range: $[-1, 1]$, higher is better
- Choose K with highest average silhouette score

Gap Statistic:

- Compare WCSS to expected WCSS under null (random) data
- Choose K where gap is largest

Domain Knowledge:

- Sometimes you know K from context
- Example: Customer segments (budget/mid/premium)
- Example: Image compression (fixed color palette)

When K-Means Fails

K-Means Assumptions

K-Means works well when:

1. Clusters are **spherical** (same in all directions)

K-Means Assumptions

K-Means works well when:

1. Clusters are **spherical** (same in all directions)
2. Clusters have **similar sizes**

K-Means Assumptions

K-Means works well when:

1. Clusters are **spherical** (same in all directions)
2. Clusters have **similar sizes**
3. Clusters have **similar densities**

K-Means Assumptions

K-Means works well when:

1. Clusters are **spherical** (same in all directions)
2. Clusters have **similar sizes**
3. Clusters have **similar densities**
4. Data is **isotropic** (no strong directional patterns)

K-Means Assumptions

K-Means works well when:

1. Clusters are **spherical** (same in all directions)
2. Clusters have **similar sizes**
3. Clusters have **similar densities**
4. Data is **isotropic** (no strong directional patterns)

Important:

Reality: Many real datasets violate these assumptions!
Let's see what happens...

Failure Mode 1: Non-Spherical Clusters

Problem: K-Means assumes clusters are "ball-shaped"

Failure Mode 1: Non-Spherical Clusters

Problem: K-Means assumes clusters are "ball-shaped"

Example: Two elongated, parallel clusters

Failure Mode 1: Non-Spherical Clusters

Problem: K-Means assumes clusters are "ball-shaped"

Example: Two elongated, parallel clusters

What happens:

- K-Means uses Euclidean distance
- Creates **spherical** decision boundaries
- Incorrectly splits elongated clusters

Failure Mode 1: Non-Spherical Clusters

Problem: K-Means assumes clusters are "ball-shaped"

Example: Two elongated, parallel clusters

What happens:

- K-Means uses Euclidean distance
- Creates **spherical** decision boundaries
- Incorrectly splits elongated clusters

Solution:

- Transform data first (e.g., PCA rotation)
- Use Gaussian Mixture Models (handles ellipsoids)
- Use spectral clustering

Failure Mode 2: Non-Convex Shapes

Problem: K-Means creates convex decision boundaries

Failure Mode 2: Non-Convex Shapes

Problem: K-Means creates convex decision boundaries

Example: Two interleaving crescents (moons)

Failure Mode 2: Non-Convex Shapes

Problem: K-Means creates convex decision boundaries

Example: Two interleaving crescents (moons)

What happens:

- Cannot separate non-convex shapes
- Centroids end up in wrong regions
- Boundaries cut through clusters

Failure Mode 2: Non-Convex Shapes

Problem: K-Means creates convex decision boundaries

Example: Two interleaving crescents (moons)

What happens:

- Cannot separate non-convex shapes
- Centroids end up in wrong regions
- Boundaries cut through clusters

Solution:

- DBSCAN (handles arbitrary shapes)
- Spectral clustering
- Kernel K-Means (kernel trick!)

Failure Mode 3: Different Cluster Sizes

Problem: K-Means prefers equal-sized clusters

Failure Mode 3: Different Cluster Sizes

Problem: K-Means prefers equal-sized clusters

Example: One large cluster + one tiny cluster

Failure Mode 3: Different Cluster Sizes

Problem: K-Means prefers equal-sized clusters

Example: One large cluster + one tiny cluster

What happens:

- Large cluster gets split into multiple parts
- Small cluster gets merged with nearby large cluster
- Objective function is minimized, but result is wrong!

Failure Mode 3: Different Cluster Sizes

Problem: K-Means prefers equal-sized clusters

Example: One large cluster + one tiny cluster

What happens:

- Large cluster gets split into multiple parts
- Small cluster gets merged with nearby large cluster
- Objective function is minimized, but result is wrong!

Why?

- Splitting large cluster reduces more WCSS
- Algorithm is greedy - doesn't know true structure

Failure Mode 3: Different Cluster Sizes

Problem: K-Means prefers equal-sized clusters

Example: One large cluster + one tiny cluster

What happens:

- Large cluster gets split into multiple parts
- Small cluster gets merged with nearby large cluster
- Objective function is minimized, but result is wrong!

Why?

- Splitting large cluster reduces more WCSS
- Algorithm is greedy - doesn't know true structure

Solution:

- Weighted K-Means
- DBSCAN (density-based)
- Hierarchical clustering

Failure Mode 4: Different Densities

Problem: K-Means assumes similar density across clusters

Failure Mode 4: Different Densities

Problem: K-Means assumes similar density across clusters

Example: Dense cluster next to sparse cluster

Failure Mode 4: Different Densities

Problem: K-Means assumes similar density across clusters

Example: Dense cluster next to sparse cluster

What happens:

- Dense cluster may be over-segmented
- Sparse cluster absorbs nearby points from dense cluster
- Boundary is placed incorrectly

Failure Mode 4: Different Densities

Problem: K-Means assumes similar density across clusters

Example: Dense cluster next to sparse cluster

What happens:

- Dense cluster may be over-segmented
- Sparse cluster absorbs nearby points from dense cluster
- Boundary is placed incorrectly

Solution:

- DBSCAN (explicitly models density)
- HDBSCAN (hierarchical density-based)
- GMM with different covariances

Failure Mode 5: Outliers

Problem: K-Means uses squared distances - very sensitive to outliers!

Failure Mode 5: Outliers

Problem: K-Means uses squared distances - very sensitive to outliers!

Example: Clean cluster + a few outliers

Failure Mode 5: Outliers

Problem: K-Means uses squared distances - very sensitive to outliers!

Example: Clean cluster + a few outliers

What happens:

- Outliers "pull" centroids away from true cluster centers
- Can create spurious clusters for noise points
- Squared distance amplifies effect ($100^2 = 10000!$)

Failure Mode 5: Outliers

Problem: K-Means uses squared distances - very sensitive to outliers!

Example: Clean cluster + a few outliers

What happens:

- Outliers "pull" centroids away from true cluster centers
- Can create spurious clusters for noise points
- Squared distance amplifies effect ($100^2 = 10000!$)

Solutions:

- Pre-process: Remove outliers first
- K-Medoids (uses medians, more robust)
- Trimmed K-Means (ignores worst $\alpha\%$ of points)
- DBSCAN (marks outliers as noise)

Failure Mode 6: Bad Initialization

Problem: Random initialization can lead to poor local minima

Failure Mode 6: Bad Initialization

Problem: Random initialization can lead to poor local minima

Example: All initial centroids in same region

Failure Mode 6: Bad Initialization

Problem: Random initialization can lead to poor local minima

Example: All initial centroids in same region

What happens:

- Converges to nearby local optimum
- Misses true clusters far away
- Different runs give different results!

Failure Mode 6: Bad Initialization

Problem: Random initialization can lead to poor local minima

Example: All initial centroids in same region

What happens:

- Converges to nearby local optimum
- Misses true clusters far away
- Different runs give different results!

Solution: K-Means++ (discussed next!)

- Smart initialization - spread out initial centroids
- Much more likely to find good solution
- Now default in most libraries

K-Means++

K-Means++: Smarter Initialization

Problem with random init: Can start with all centroids clustered together

K-Means++: Smarter Initialization

Problem with random init: Can start with all centroids clustered together

K-Means++ Idea: Choose initial centroids that are **far apart**

K-Means++: Smarter Initialization

Problem with random init: Can start with all centroids clustered together

K-Means++ Idea: Choose initial centroids that are **far apart**
Key Insight:

- If true clusters are far apart
- Initial centroids should also be far apart
- More likely to have one centroid per true cluster

K-Means++: Smarter Initialization

Problem with random init: Can start with all centroids clustered together

K-Means++ Idea: Choose initial centroids that are **far apart**
Key Insight:

- If true clusters are far apart
- Initial centroids should also be far apart
- More likely to have one centroid per true cluster

Key Points: Key Points

How to choose "far apart" centroids?

Use **weighted random sampling** - points far from existing centroids have higher probability!

K-Means++ Algorithm

Input: Data points $\{x_1, \dots, x_n\}$, number of clusters K

K-Means++ Algorithm

Input: Data points $\{x_1, \dots, x_n\}$, number of clusters K

Initialization:

1. Choose first centroid μ_1 uniformly at random from data

K-Means++ Algorithm

Input: Data points $\{x_1, \dots, x_n\}$, number of clusters K

Initialization:

1. Choose first centroid μ_1 uniformly at random from data
2. For $k = 2, 3, \dots, K$:

K-Means++ Algorithm

Input: Data points $\{x_1, \dots, x_n\}$, number of clusters K

Initialization:

1. Choose first centroid μ_1 uniformly at random from data
2. For $k = 2, 3, \dots, K$:
 - 1) For each point x_i , compute:

$$D(x_i) = \min_{j < k} \|x_i - \mu_j\|^2$$

(squared distance to nearest already-chosen centroid)

K-Means++ Algorithm

Input: Data points $\{x_1, \dots, x_n\}$, number of clusters K

Initialization:

1. Choose first centroid μ_1 uniformly at random from data
2. For $k = 2, 3, \dots, K$:
 - 1) For each point x_i , compute:

$$D(x_i) = \min_{j < k} \|x_i - \mu_j\|^2$$

(squared distance to nearest already-chosen centroid)

- 2) Choose next centroid μ_k with probability:

$$P(x_i) = \frac{D(x_i)}{\sum_{j=1}^n D(x_j)}$$

K-Means++ Algorithm

Input: Data points $\{x_1, \dots, x_n\}$, number of clusters K

Initialization:

1. Choose first centroid μ_1 uniformly at random from data
2. For $k = 2, 3, \dots, K$:
 - 1) For each point x_i , compute:

$$D(x_i) = \min_{j < k} \|x_i - \mu_j\|^2$$

(squared distance to nearest already-chosen centroid)

- 2) Choose next centroid μ_k with probability:

$$P(x_i) = \frac{D(x_i)}{\sum_{j=1}^n D(x_j)}$$

Then: Run standard K-Means with these initial centroids

K-Means++: Why It Works

Intuition:

- Points **far** from existing centroids have high $D(x_i)^2$

K-Means++: Why It Works

Intuition:

- Points **far** from existing centroids have high $D(x_i)^2$
- Points **near** existing centroids have low $D(x_i)^2$

K-Means++: Why It Works

Intuition:

- Points **far** from existing centroids have high $D(x_i)^2$
- Points **near** existing centroids have low $D(x_i)^2$
- Weighted sampling favors distant points

K-Means++: Why It Works

Intuition:

- Points **far** from existing centroids have high $D(x_i)^2$
- Points **near** existing centroids have low $D(x_i)^2$
- Weighted sampling favors distant points
- Result: Centroids spread out across data

K-Means++: Why It Works

Intuition:

- Points **far** from existing centroids have high $D(x_i)^2$
- Points **near** existing centroids have low $D(x_i)^2$
- Weighted sampling favors distant points
- Result: Centroids spread out across data

Example: Concrete Example

Suppose we've chosen $\mu_1 = (0, 0)$. Consider two points:

- $x_a = (1, 0)$: $D(x_a) = 1 \rightarrow P(x_a) \propto 1$
- $x_b = (10, 0)$: $D(x_b) = 100 \rightarrow P(x_b) \propto 100$

Point x_b is $100\times$ more likely to be chosen as μ_2 !

K-Means++: Theoretical Guarantee

Theorem (Arthur & Vassilvitskii, 2007):

K-Means++ initialization gives an expected approximation ratio of $O(\log K)$ to the optimal clustering.

K-Means++: Theoretical Guarantee

Theorem (Arthur & Vassilvitskii, 2007):

K-Means++ initialization gives an expected approximation ratio of $O(\log K)$ to the optimal clustering.

In English:

- Let $\text{WCSS}_{\text{optimal}} = \text{best possible WCSS}$
- Let $\text{WCSS}_{\text{KM}++} = \text{WCSS from K-Means++ (after convergence)}$
- Then:

$$\mathbb{E}[\text{WCSS}_{\text{KM}++}] \leq O(\log K) \cdot \text{WCSS}_{\text{optimal}}$$

K-Means++: Theoretical Guarantee

Theorem (Arthur & Vassilvitskii, 2007):

K-Means++ initialization gives an expected approximation ratio of $O(\log K)$ to the optimal clustering.

In English:

- Let $\text{WCSS}_{\text{optimal}} = \text{best possible WCSS}$
- Let $\text{WCSS}_{\text{KM}++} = \text{WCSS from K-Means++ (after convergence)}$
- Then:

$$\mathbb{E}[\text{WCSS}_{\text{KM}++}] \leq O(\log K) \cdot \text{WCSS}_{\text{optimal}}$$

Random initialization: No such guarantee!

K-Means++: Theoretical Guarantee

Theorem (Arthur & Vassilvitskii, 2007):

K-Means++ initialization gives an expected approximation ratio of $O(\log K)$ to the optimal clustering.

In English:

- Let $\text{WCSS}_{\text{optimal}} = \text{best possible WCSS}$
- Let $\text{WCSS}_{\text{KM}++} = \text{WCSS from K-Means++ (after convergence)}$
- Then:

$$\mathbb{E}[\text{WCSS}_{\text{KM}++}] \leq O(\log K) \cdot \text{WCSS}_{\text{optimal}}$$

Random initialization: No such guarantee!

Key Points: Key Points

K-Means++ is **provably better** than random initialization
This is why it's now the default in scikit-learn, etc.

K-Means++ vs Standard K-Means

Property	Random Init	K-Means++
Convergence guarantee	Yes	Yes
Quality guarantee	No	$O(\log K)$ -approx
Iterations needed	More	Fewer
Consistency across runs	Poor	Better
Initialization time	$O(1)$	$O(nKd)$
Total time	$O(nKdt)$	$O(nKd(t + K))$

K-Means++ vs Standard K-Means

Property	Random Init	K-Means++
Convergence guarantee	Yes	Yes
Quality guarantee	No	$O(\log K)$ -approx
Iterations needed	More	Fewer
Consistency across runs	Poor	Better
Initialization time	$O(1)$	$O(nKd)$
Total time	$O(nKdt)$	$O(nKd(t + K))$

Practical Advice:

- Always use K-Means++ (default in scikit-learn)
- Initialization cost is negligible compared to iterations
- Better results with fewer iterations

Hierarchical Clustering

Limitations of K-Means

K-Means requires:

- Specifying K beforehand
- Running multiple times with different K (elbow method)
- Assumes spherical clusters

Limitations of K-Means

K-Means requires:

- Specifying K beforehand
- Running multiple times with different K (elbow method)
- Assumes spherical clusters

Can we do better?

- Get clustering at **all levels** of granularity
- From $K = n$ (each point alone) to $K = 1$ (everything together)
- Visualize as a tree (dendrogram)
- Cut tree at any height to get desired K

Limitations of K-Means

K-Means requires:

- Specifying K beforehand
- Running multiple times with different K (elbow method)
- Assumes spherical clusters

Can we do better?

- Get clustering at **all levels** of granularity
- From $K = n$ (each point alone) to $K = 1$ (everything together)
- Visualize as a tree (dendrogram)
- Cut tree at any height to get desired K

Key Points: Key Points

Hierarchical Clustering: Build a tree of nested clusters!

Two Approaches to Hierarchical Clustering

1. Agglomerative (Bottom-Up)

- Start: n clusters (each point alone)
- Repeat: Merge two closest clusters
- End: 1 cluster (everything together)
- Most common approach

Two Approaches to Hierarchical Clustering

1. Agglomerative (Bottom-Up)

- Start: n clusters (each point alone)
- Repeat: Merge two closest clusters
- End: 1 cluster (everything together)
- Most common approach

Intuition: Like building from atoms to molecules to structures

Two Approaches to Hierarchical Clustering

1. Agglomerative (Bottom-Up)

- Start: n clusters (each point alone)
- Repeat: Merge two closest clusters
- End: 1 cluster (everything together)
- **Most common approach**

2. Divisive (Top-Down)

- Start: 1 cluster (everything together)
- Repeat: Split cluster into two
- End: n clusters (each point alone)
- Less common
(computationally harder)

Intuition: Like building from atoms to molecules to structures

Two Approaches to Hierarchical Clustering

1. Agglomerative (Bottom-Up)

- Start: n clusters (each point alone)
- Repeat: Merge two closest clusters
- End: 1 cluster (everything together)
- **Most common approach**

Intuition: Like building from atoms to molecules to structures

2. Divisive (Top-Down)

- Start: 1 cluster (everything together)
- Repeat: Split cluster into two
- End: n clusters (each point alone)
- Less common
(computationally harder)

Intuition: Like breaking a rock into smaller and smaller pieces

Two Approaches to Hierarchical Clustering

1. Agglomerative (Bottom-Up)

- Start: n clusters (each point alone)
- Repeat: Merge two closest clusters
- End: 1 cluster (everything together)
- **Most common approach**

Intuition: Like building from atoms to molecules to structures

2. Divisive (Top-Down)

- Start: 1 cluster (everything together)
- Repeat: Split cluster into two
- End: n clusters (each point alone)
- Less common (computationally harder)

Intuition: Like breaking a rock into smaller and smaller pieces

Important:

Agglomerative Clustering: Algorithm

Input: Data points $\{x_1, \dots, x_n\}$

Algorithm:

1. **Initialize:** Start with n clusters, one per point

$$C_i = \{x_i\} \quad \text{for } i = 1, \dots, n$$

Agglomerative Clustering: Algorithm

Input: Data points $\{x_1, \dots, x_n\}$

Algorithm:

1. **Initialize:** Start with n clusters, one per point

$$C_i = \{x_i\} \quad \text{for } i = 1, \dots, n$$

2. **Repeat** until only 1 cluster remains:

Agglomerative Clustering: Algorithm

Input: Data points $\{x_1, \dots, x_n\}$

Algorithm:

1. **Initialize:** Start with n clusters, one per point

$$C_i = \{x_i\} \quad \text{for } i = 1, \dots, n$$

2. **Repeat** until only 1 cluster remains:

- 1) Find the two "closest" clusters C_i and C_j

Agglomerative Clustering: Algorithm

Input: Data points $\{x_1, \dots, x_n\}$

Algorithm:

1. **Initialize:** Start with n clusters, one per point

$$C_i = \{x_i\} \quad \text{for } i = 1, \dots, n$$

2. **Repeat** until only 1 cluster remains:

- 1) Find the two "closest" clusters C_i and C_j
- 2) Merge them: $C_{\text{new}} = C_i \cup C_j$

Agglomerative Clustering: Algorithm

Input: Data points $\{x_1, \dots, x_n\}$

Algorithm:

1. **Initialize:** Start with n clusters, one per point

$$C_i = \{x_i\} \quad \text{for } i = 1, \dots, n$$

2. **Repeat** until only 1 cluster remains:

- 1) Find the two "closest" clusters C_i and C_j
- 2) Merge them: $C_{\text{new}} = C_i \cup C_j$
- 3) Record merge in tree structure (dendrogram)

Agglomerative Clustering: Algorithm

Input: Data points $\{x_1, \dots, x_n\}$

Algorithm:

1. **Initialize:** Start with n clusters, one per point

$$C_i = \{x_i\} \quad \text{for } i = 1, \dots, n$$

2. **Repeat** until only 1 cluster remains:

- 1) Find the two "closest" clusters C_i and C_j
- 2) Merge them: $C_{\text{new}} = C_i \cup C_j$
- 3) Record merge in tree structure (dendrogram)

3. **Output:** Dendrogram (tree showing all merges)

Agglomerative Clustering: Algorithm

Input: Data points $\{x_1, \dots, x_n\}$

Algorithm:

1. **Initialize:** Start with n clusters, one per point

$$C_i = \{x_i\} \quad \text{for } i = 1, \dots, n$$

2. **Repeat** until only 1 cluster remains:

- 1) Find the two "closest" clusters C_i and C_j
- 2) Merge them: $C_{\text{new}} = C_i \cup C_j$
- 3) Record merge in tree structure (dendrogram)

3. **Output:** Dendrogram (tree showing all merges)

Key Points: Key Points

Key Question: How do we measure distance between **clusters** (not just points)?

Linkage Criteria: Measuring Cluster Distance

Given two clusters C_i and C_j , what is their distance?

Linkage Criteria: Measuring Cluster Distance

Given two clusters C_i and C_j , what is their distance?

1. Single Linkage (Minimum):

$$d(C_i, C_j) = \min_{x \in C_i, y \in C_j} \|x - y\|$$

- Distance between closest points
- Can create "chains" (long, thin clusters)

Linkage Criteria: Measuring Cluster Distance

Given two clusters C_i and C_j , what is their distance?

1. Single Linkage (Minimum):

$$d(C_i, C_j) = \min_{x \in C_i, y \in C_j} \|x - y\|$$

- Distance between closest points
- Can create "chains" (long, thin clusters)

2. Complete Linkage (Maximum):

$$d(C_i, C_j) = \max_{x \in C_i, y \in C_j} \|x - y\|$$

- Distance between farthest points
- Creates compact, spherical clusters

Linkage Criteria: Measuring Cluster Distance

Given two clusters C_i and C_j , what is their distance?

1. Single Linkage (Minimum):

$$d(C_i, C_j) = \min_{x \in C_i, y \in C_j} \|x - y\|$$

- Distance between closest points
- Can create "chains" (long, thin clusters)

2. Complete Linkage (Maximum):

$$d(C_i, C_j) = \max_{x \in C_i, y \in C_j} \|x - y\|$$

- Distance between farthest points
- Creates compact, spherical clusters

3. Average Linkage:

$$d(C_i, C_j) = \frac{1}{|C_i||C_j|} \sum_i \sum_j \|x - y\|$$

More Linkage Criteria

4. Centroid Linkage:

$$d(C_i, C_j) = \|\mu_i - \mu_j\|$$

where $\mu_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$

- Distance between centroids
- Can have inversions (not monotonic!)

More Linkage Criteria

4. Centroid Linkage:

$$d(C_i, C_j) = \|\mu_i - \mu_j\|$$

where $\mu_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$

- Distance between centroids
- Can have inversions (not monotonic!)

5. Ward's Method (Minimum Variance):

$$d(C_i, C_j) = \frac{|C_i||C_j|}{|C_i| + |C_j|} \|\mu_i - \mu_j\|^2$$

- Minimizes increase in total within-cluster variance
- Tends to create equal-sized clusters
- **Most popular choice** in practice

More Linkage Criteria

4. Centroid Linkage:

$$d(C_i, C_j) = \|\mu_i - \mu_j\|$$

where $\mu_i = \frac{1}{|C_i|} \sum_{x \in C_i} x$

- Distance between centroids
- Can have inversions (not monotonic!)

5. Ward's Method (Minimum Variance):

$$d(C_i, C_j) = \frac{|C_i||C_j|}{|C_i| + |C_j|} \|\mu_i - \mu_j\|^2$$

- Minimizes increase in total within-cluster variance
- Tends to create equal-sized clusters
- **Most popular choice** in practice

Key Points: Key Points

Dendrogram: Visualizing the Hierarchy

Dendrogram = Tree diagram showing merge history

Dendrogram: Visualizing the Hierarchy

Dendrogram = Tree diagram showing merge history

Reading a dendrogram:

- **Bottom:** Individual points
- **Height:** Distance at which clusters merge
- **Horizontal cut:** Determines number of clusters

Dendrogram: Visualizing the Hierarchy

Dendrogram = Tree diagram showing merge history

Reading a dendrogram:

- **Bottom:** Individual points
- **Height:** Distance at which clusters merge
- **Horizontal cut:** Determines number of clusters

Example:

- Cut at height $h = 2$: Might give 5 clusters
- Cut at height $h = 5$: Might give 2 clusters
- Cut at height $h = 10$: Everything in 1 cluster

Dendrogram: Visualizing the Hierarchy

Dendrogram = Tree diagram showing merge history

Reading a dendrogram:

- **Bottom:** Individual points
- **Height:** Distance at which clusters merge
- **Horizontal cut:** Determines number of clusters

Example:

- Cut at height $h = 2$: Might give 5 clusters
- Cut at height $h = 5$: Might give 2 clusters
- Cut at height $h = 10$: Everything in 1 cluster

Example: Flexibility

Advantage: Don't need to specify K beforehand!

Run once, then choose K by cutting dendrogram at desired height.

Hierarchical Clustering: Complexity

Naive Algorithm:

- Need to merge $n - 1$ times
- At each step, compute distances between all cluster pairs
- Worst case: $O(n^3)$ time

Hierarchical Clustering: Complexity

Naive Algorithm:

- Need to merge $n - 1$ times
- At each step, compute distances between all cluster pairs
- Worst case: $O(n^3)$ time

Optimized (using priority queue):

- Time: $O(n^2 \log n)$
- Space: $O(n^2)$ (store pairwise distances)

Hierarchical Clustering: Complexity

Naive Algorithm:

- Need to merge $n - 1$ times
- At each step, compute distances between all cluster pairs
- Worst case: $O(n^3)$ time

Optimized (using priority queue):

- Time: $O(n^2 \log n)$
- Space: $O(n^2)$ (store pairwise distances)

Comparison to K-Means:

	K-Means	Hierarchical
Time	$O(nKdt)$	$O(n^2 \log n)$
Space	$O((n + K)d)$	$O(n^2)$

Hierarchical Clustering: Complexity

Naive Algorithm:

- Need to merge $n - 1$ times
- At each step, compute distances between all cluster pairs
- Worst case: $O(n^3)$ time

Optimized (using priority queue):

- Time: $O(n^2 \log n)$
- Space: $O(n^2)$ (store pairwise distances)

Comparison to K-Means:

	K-Means	Hierarchical
Time	$O(nKdt)$	$O(n^2 \log n)$
Space	$O((n + K)d)$	$O(n^2)$

Important: H

ierarchical is **much slower** for large n (thousands+)

K-Means scales better to large datasets

Hierarchical vs K-Means: When to Use Each?

Property	K-Means	Hierarchical
Specify K ?	Yes	No
Speed (large n)	Fast	Slow
Deterministic?	No	Yes
Cluster shape	Spherical	Flexible
Dendrogram?	No	Yes
Scalability	Excellent	Poor

Hierarchical vs K-Means: When to Use Each?

Property	K-Means	Hierarchical
Specify K ?	Yes	No
Speed (large n)	Fast	Slow
Deterministic?	No	Yes
Cluster shape	Spherical	Flexible
Dendrogram?	No	Yes
Scalability	Excellent	Poor

Use K-Means when:

- Large dataset ($n > 10000$)
- Roughly spherical clusters
- Know K (or can use elbow method)

Hierarchical vs K-Means: When to Use Each?

Property	K-Means	Hierarchical
Specify K ?	Yes	No
Speed (large n)	Fast	Slow
Deterministic?	No	Yes
Cluster shape	Spherical	Flexible
Dendrogram?	No	Yes
Scalability	Excellent	Poor

Use K-Means when:

- Large dataset ($n > 10000$)
- Roughly spherical clusters
- Know K (or can use elbow method)

Use Hierarchical when:

- Small-medium dataset ($n < 5000$)
- Want full hierarchy (dendrogram)
- Don't know K beforehand

DBSCAN

Limitations of K-Means and Hierarchical

Both algorithms struggle with:

- Non-convex shapes (crescents, rings)

Limitations of K-Means and Hierarchical

Both algorithms struggle with:

- Non-convex shapes (crescents, rings)
- Outliers and noise

Limitations of K-Means and Hierarchical

Both algorithms struggle with:

- Non-convex shapes (crescents, rings)
- Outliers and noise
- Varying density clusters

Limitations of K-Means and Hierarchical

Both algorithms struggle with:

- Non-convex shapes (crescents, rings)
- Outliers and noise
- Varying density clusters

Example: Motivating Example

Consider two clusters:

- Dense core of points (100 points in small area)
- Scattered points around core (10 points far away)

Question: Should scattered points be their own clusters or noise?

Limitations of K-Means and Hierarchical

Both algorithms struggle with:

- Non-convex shapes (crescents, rings)
- Outliers and noise
- Varying density clusters

Example: Motivating Example

Consider two clusters:

- Dense core of points (100 points in small area)
- Scattered points around core (10 points far away)

Question: Should scattered points be their own clusters or noise?

DBSCAN: Density-Based Spatial Clustering

Core Idea:

A cluster is a region where points are densely packed together

DBSCAN: Density-Based Spatial Clustering

Core Idea:

A cluster is a region where points are densely packed together

Advantages:

- Can find **arbitrarily shaped** clusters
- Robust to **outliers** (marks them as noise)
- No need to specify number of clusters K
- Deterministic (same result every time)

DBSCAN: Density-Based Spatial Clustering

Core Idea:

A cluster is a region where points are densely packed together

Advantages:

- Can find **arbitrarily shaped** clusters
- Robust to **outliers** (marks them as noise)
- No need to specify number of clusters K
- Deterministic (same result every time)

Parameters:

- ε (epsilon): Radius of neighborhood
- MinPts: Minimum number of points in neighborhood to be "dense"

DBSCAN: Key Definitions

1. **ε -neighborhood** of point p :

$$N_\varepsilon(p) = \{q : \|p - q\| \leq \varepsilon\}$$

All points within radius ε of p

DBSCAN: Key Definitions

1. **ε -neighborhood** of point p :

$$N_\varepsilon(p) = \{q : \|p - q\| \leq \varepsilon\}$$

All points within radius ε of p

2. **Core point**:

- Point p is a **core point** if $|N_\varepsilon(p)| \geq \text{MinPts}$
- Has at least MinPts neighbors (including itself)
- These form the "dense regions"

DBSCAN: Key Definitions

1. ε -neighborhood of point p :

$$N_\varepsilon(p) = \{q : \|p - q\| \leq \varepsilon\}$$

All points within radius ε of p

2. Core point:

- Point p is a **core point** if $|N_\varepsilon(p)| \geq \text{MinPts}$
- Has at least MinPts neighbors (including itself)
- These form the "dense regions"

3. Border point:

- Not a core point
- But is in $N_\varepsilon(p)$ of some core point p
- On the "edge" of a cluster

DBSCAN: Key Definitions

1. ε -neighborhood of point p :

$$N_\varepsilon(p) = \{q : \|p - q\| \leq \varepsilon\}$$

All points within radius ε of p

2. Core point:

- Point p is a **core point** if $|N_\varepsilon(p)| \geq \text{MinPts}$
- Has at least MinPts neighbors (including itself)
- These form the "dense regions"

3. Border point:

- Not a core point
- But is in $N_\varepsilon(p)$ of some core point p
- On the "edge" of a cluster

4. Noise point:

- Neither core nor border
- Isolated outliers

DBSCAN: Visual Example of Point Types

Setup: $\varepsilon = 1$, MinPts = 4

[Visual would show: points with circles of radius ε]

DBSCAN: Visual Example of Point Types

Setup: $\varepsilon = 1$, MinPts = 4

[Visual would show: points with circles of radius ε]

Point Classification:

- **Point A:** Has 5 neighbors → **Core point**
- **Point B:** Has 6 neighbors → **Core point**
- **Point C:** Has 2 neighbors, but in neighborhood of A → **Border point**
- **Point D:** Has 1 neighbor, not near any core → **Noise**

DBSCAN: Visual Example of Point Types

Setup: $\varepsilon = 1$, MinPts = 4

[Visual would show: points with circles of radius ε]

Point Classification:

- **Point A:** Has 5 neighbors → **Core point**
- **Point B:** Has 6 neighbors → **Core point**
- **Point C:** Has 2 neighbors, but in neighborhood of A → **Border point**
- **Point D:** Has 1 neighbor, not near any core → **Noise**

Key Points: Key Points

Cluster = All core points connected + their border points

DBSCAN: Density Connectivity

Direct density-reachable:

- Point q is directly density-reachable from p if:
 1. p is a core point
 2. $q \in N_\varepsilon(p)$

DBSCAN: Density Connectivity

Direct density-reachable:

- Point q is directly density-reachable from p if:
 1. p is a core point
 2. $q \in N_\varepsilon(p)$

Density-reachable:

- q is density-reachable from p if there's a chain:

$$p = p_1, p_2, \dots, p_n = q$$

where each p_{i+1} is directly density-reachable from p_i

DBSCAN: Density Connectivity

Direct density-reachable:

- Point q is directly density-reachable from p if:
 1. p is a core point
 2. $q \in N_\epsilon(p)$

Density-reachable:

- q is density-reachable from p if there's a chain:

$$p = p_1, p_2, \dots, p_n = q$$

where each p_{i+1} is directly density-reachable from p_i

Density-connected:

- Points p and q are density-connected if both are density-reachable from some point o

DBSCAN: Density Connectivity

Direct density-reachable:

- Point q is directly density-reachable from p if:
 1. p is a core point
 2. $q \in N_\epsilon(p)$

Density-reachable:

- q is density-reachable from p if there's a chain:

$$p = p_1, p_2, \dots, p_n = q$$

where each p_{i+1} is directly density-reachable from p_i

Density-connected:

- Points p and q are density-connected if both are density-reachable from some point o

Definition (DBSCAN Cluster)

A **cluster** is the maximal set of density-connected points.

DBSCAN Algorithm

Input: Data points, ε , MinPts

Algorithm:

1. Mark all points as UNVISITED

DBSCAN Algorithm

Input: Data points, ε , MinPts

Algorithm:

1. Mark all points as UNVISITED
2. **For each** unvisited point p :

DBSCAN Algorithm

Input: Data points, ε , MinPts

Algorithm:

1. Mark all points as UNVISITED
2. **For each** unvisited point p :
 - 1) Mark p as VISITED

DBSCAN Algorithm

Input: Data points, ε , MinPts

Algorithm:

1. Mark all points as UNVISITED
2. **For each** unvisited point p :
 - 1) Mark p as VISITED
 - 2) Find $N_\varepsilon(p)$ (all neighbors within ε)

DBSCAN Algorithm

Input: Data points, ε , MinPts

Algorithm:

1. Mark all points as UNVISITED
2. **For each** unvisited point p :
 - 1) Mark p as VISITED
 - 2) Find $N_\varepsilon(p)$ (all neighbors within ε)
 - 3) **If** $|N_\varepsilon(p)| < \text{MinPts}$:
 - Mark p as NOISE (may change later!)

DBSCAN Algorithm

Input: Data points, ε , MinPts

Algorithm:

1. Mark all points as UNVISITED
2. **For each** unvisited point p :
 - 1) Mark p as VISITED
 - 2) Find $N_\varepsilon(p)$ (all neighbors within ε)
 - 3) **If** $|N_\varepsilon(p)| < \text{MinPts}$:
 - Mark p as NOISE (may change later!)
 - 4) **Else** (p is core point):
 - Create new cluster C
 - Add p to C
 - Expand cluster: recursively add all density-reachable points

DBSCAN Algorithm

Input: Data points, ε , MinPts

Algorithm:

1. Mark all points as UNVISITED
2. **For each** unvisited point p :
 - 1) Mark p as VISITED
 - 2) Find $N_\varepsilon(p)$ (all neighbors within ε)
 - 3) **If** $|N_\varepsilon(p)| < \text{MinPts}$:
 - Mark p as NOISE (may change later!)
 - 4) **Else** (p is core point):
 - Create new cluster C
 - Add p to C
 - Expand cluster: recursively add all density-reachable points
3. **Output:** Set of clusters + noise points

DBSCAN: Expand Cluster Procedure

ExpandCluster($p, N_\varepsilon(p), C$):

1. Add p to cluster C
2. **For each** point $q \in N_\varepsilon(p)$:
 - 1) **If** q is UNVISITED:
 - Mark q as VISITED
 - Find $N_\varepsilon(q)$
 - **If** $|N_\varepsilon(q)| \geq \text{MinPts}$:
 - q is a core point!
 - Add $N_\varepsilon(q)$ to queue (will expand recursively)
 - 2) **If** q not yet in any cluster:
 - Add q to C

DBSCAN: Expand Cluster Procedure

ExpandCluster($p, N_\varepsilon(p), C$):

1. Add p to cluster C
2. **For each** point $q \in N_\varepsilon(p)$:
 - 1) **If** q is UNVISITED:
 - Mark q as VISITED
 - Find $N_\varepsilon(q)$
 - **If** $|N_\varepsilon(q)| \geq \text{MinPts}$:
 - q is a core point!
 - Add $N_\varepsilon(q)$ to queue (will expand recursively)
 - 2) **If** q not yet in any cluster:
 - Add q to C

Key Points: Key Points

Result: All density-reachable points from p are in cluster C

DBSCAN: Worked Example (1/4)

Setup: $\varepsilon = 1.5$, MinPts = 3

Data: 10 points in 2D

Point	Coords	Point	Coords
A	(1, 1)	F	(2, 2)
B	(1, 2)	G	(10, 10)
C	(2, 1)	H	(10, 11)
D	(1.5, 1.5)	I	(11, 10)
E	(2, 3)	J	(5, 5)

Visualize: Three regions - dense cluster left, dense cluster right, one isolated point

DBSCAN: Worked Example (2/4) - Finding Core Points

Check each point's neighborhood:

- **Point A** (1, 1): Neighbors = {A, B, C, D} (4 points) → **Core**
- **Point B** (1, 2): Neighbors = {A, B, D, E, F} (5 points) → **Core**
- **Point C** (2, 1): Neighbors = {A, C, D, F} (4 points) → **Core**
- **Point D** (1.5, 1.5): Neighbors = {A, B, C, D, F} (5 points) → **Core**
- **Point E** (2, 3): Neighbors = {B, E, F} (3 points) → **Core**
- **Point F** (2, 2): Neighbors = {B, C, D, E, F} (5 points) → **Core**
- **Point G** (10, 10): Neighbors = {G, H, I} (3 points) → **Core**
- **Point H** (10, 11): Neighbors = {G, H, I} (3 points) → **Core**
- **Point I** (11, 10): Neighbors = {G, H, I} (3 points) → **Core**
- **Point J** (5, 5): Neighbors = {J} (1 point) → **Noise**

DBSCAN: Worked Example (3/4) - Forming Clusters

Start with point A:

- A is core → Create **Cluster 1**
- Expand: A can reach B, C, D (all core)
- B can reach E, F (both core)
- Continue expanding: All of {A, B, C, D, E, F} are density-connected
- **Cluster 1** = {A, B, C, D, E, F}

DBSCAN: Worked Example (3/4) - Forming Clusters

Start with point A:

- A is core → Create **Cluster 1**
- Expand: A can reach B, C, D (all core)
- B can reach E, F (both core)
- Continue expanding: All of {A, B, C, D, E, F} are density-connected
- **Cluster 1** = {A, B, C, D, E, F}

Continue with unvisited point G:

- G is core → Create **Cluster 2**
- Expand: G can reach H, I (both core)
- All of {G, H, I} are density-connected
- **Cluster 2** = {G, H, I}

DBSCAN: Worked Example (3/4) - Forming Clusters

Start with point A:

- A is core → Create **Cluster 1**
- Expand: A can reach B, C, D (all core)
- B can reach E, F (both core)
- Continue expanding: All of {A, B, C, D, E, F} are density-connected
- **Cluster 1** = {A, B, C, D, E, F}

Continue with unvisited point G:

- G is core → Create **Cluster 2**
- Expand: G can reach H, I (both core)
- All of {G, H, I} are density-connected
- **Cluster 2** = {G, H, I}

DBSCAN: Worked Example (4/4) - Final Result

Final Clustering:

- **Cluster 1:** {A, B, C, D, E, F} (6 points)
- **Cluster 2:** {G, H, I} (3 points)
- **Noise:** {J} (1 point)

DBSCAN: Worked Example (4/4) - Final Result

Final Clustering:

- **Cluster 1:** {A, B, C, D, E, F} (6 points)
- **Cluster 2:** {G, H, I} (3 points)
- **Noise:** {J} (1 point)

Key Points: Key Points

Key Observations:

- Found 2 clusters automatically (no K specified!)
- Identified outlier J as noise
- Clusters can have different sizes (6 vs 3)
- Clusters based on density, not distance to centroid

DBSCAN: Worked Example (4/4) - Final Result

Final Clustering:

- **Cluster 1:** {A, B, C, D, E, F} (6 points)
- **Cluster 2:** {G, H, I} (3 points)
- **Noise:** {J} (1 point)

Key Points: Key Points

Key Observations:

- Found 2 clusters automatically (no K specified!)
- Identified outlier J as noise
- Clusters can have different sizes (6 vs 3)
- Clusters based on density, not distance to centroid

DBSCAN: Advantages

1. **Arbitrary shapes:** Can find non-convex, elongated clusters
2. **No K needed:** Number of clusters determined automatically
3. **Noise detection:** Explicitly identifies outliers
4. **Deterministic:** Same result every run
5. **Only 2 parameters:** ε and MinPts (vs K for K-Means)

DBSCAN: Advantages

1. **Arbitrary shapes:** Can find non-convex, elongated clusters
2. **No K needed:** Number of clusters determined automatically
3. **Noise detection:** Explicitly identifies outliers
4. **Deterministic:** Same result every run
5. **Only 2 parameters:** ε and MinPts (vs K for K-Means)

Example Applications:

- Geographic data (cities, disease outbreaks)
- Anomaly detection (fraud, network intrusion)
- Image segmentation (arbitrary region shapes)
- Point cloud processing (3D scanning)

DBSCAN: Disadvantages

1. **Parameter sensitivity:** Results depend heavily on ε and MinPts
2. **Varying density:** Struggles when clusters have different densities
3. **High dimensions:** "Curse of dimensionality" - distances become meaningless
4. **Complexity:** $O(n^2)$ naive, $O(n \log n)$ with spatial index

DBSCAN: Disadvantages

1. **Parameter sensitivity:** Results depend heavily on ε and MinPts
2. **Varying density:** Struggles when clusters have different densities
3. **High dimensions:** "Curse of dimensionality" - distances become meaningless
4. **Complexity:** $O(n^2)$ naive, $O(n \log n)$ with spatial index

Example of varying density problem:

- Dense cluster: 100 points in radius 1
- Sparse cluster: 10 points in radius 5
- No single ε works well for both!

DBSCAN: Disadvantages

1. **Parameter sensitivity:** Results depend heavily on ε and MinPts
2. **Varying density:** Struggles when clusters have different densities
3. **High dimensions:** "Curse of dimensionality" - distances become meaningless
4. **Complexity:** $O(n^2)$ naive, $O(n \log n)$ with spatial index

Example of varying density problem:

- Dense cluster: 100 points in radius 1
- Sparse cluster: 10 points in radius 5
- No single ε works well for both!

Solutions:

- HDBSCAN (hierarchical DBSCAN) - handles varying densities

Choosing DBSCAN Parameters

How to choose ε ?

K-distance graph method:

1. For each point, compute distance to k -th nearest neighbor (use $k = \text{MinPts}$)
2. Sort these distances
3. Plot sorted k -distances
4. Look for "elbow" - sharp increase indicates noise threshold
5. Set $\varepsilon = \text{distance at elbow}$

Choosing DBSCAN Parameters

How to choose ε ?

K-distance graph method:

1. For each point, compute distance to k -th nearest neighbor (use $k = \text{MinPts}$)
2. Sort these distances
3. Plot sorted k -distances
4. Look for "elbow" - sharp increase indicates noise threshold
5. Set $\varepsilon = \text{distance at elbow}$

How to choose MinPts?

- Rule of thumb: $\text{MinPts} \geq d + 1$ where d is dimensionality
- For 2D data: $\text{MinPts} = 3$ or 4
- Larger MinPts: More robust to noise, but may merge clusters

Choosing DBSCAN Parameters

How to choose ε ?

K-distance graph method:

1. For each point, compute distance to k -th nearest neighbor (use $k = \text{MinPts}$)
2. Sort these distances
3. Plot sorted k -distances
4. Look for "elbow" - sharp increase indicates noise threshold
5. Set $\varepsilon = \text{distance at elbow}$

How to choose MinPts?

- Rule of thumb: $\text{MinPts} \geq d + 1$ where d is dimensionality
- For 2D data: $\text{MinPts} = 3$ or 4
- Larger MinPts: More robust to noise, but may merge clusters

Comparing Clustering Methods

Clustering Algorithms: Summary Table

Algorithm	Shape	K?	Outliers	Complexity
K-Means	Spherical	Yes	Sensitive	$O(nKdt)$
K-Means++	Spherical	Yes	Sensitive	$O(nKd(t + K))$
Hierarchical	Flexible	No	Sensitive	$O(n^2 \log n)$
DBSCAN	Arbitrary	No	Robust	$O(n \log n)^*$

*With spatial indexing; $O(n^2)$ naive

Clustering Algorithms: Summary Table

Algorithm	Shape	K?	Outliers	Complexity
K-Means	Spherical	Yes	Sensitive	$O(nKdt)$
K-Means++	Spherical	Yes	Sensitive	$O(nKd(t + K))$
Hierarchical	Flexible	No	Sensitive	$O(n^2 \log n)$
DBSCAN	Arbitrary	No	Robust	$O(n \log n)^*$

*With spatial indexing; $O(n^2)$ naive

Additional Properties:

Algorithm	Deterministic	Scalability	High-D OK?
K-Means	No	Excellent	Yes
K-Means++	No	Excellent	Yes
Hierarchical	Yes	Poor	Moderate
DBSCAN	Yes	Good	No

Decision Tree: Which Algorithm to Use?

Choosing a Clustering Algorithm

1. Do you know K ?

- Yes → Consider K-Means++
- No → Consider Hierarchical or DBSCAN

Decision Tree: Which Algorithm to Use?

Choosing a Clustering Algorithm

1. Do you know K ?

- Yes → Consider K-Means++
- No → Consider Hierarchical or DBSCAN

2. Are clusters roughly spherical?

- Yes → K-Means++ is great
- No → DBSCAN or Hierarchical

Decision Tree: Which Algorithm to Use?

Choosing a Clustering Algorithm

1. Do you know K ?

- Yes → Consider K-Means++
- No → Consider Hierarchical or DBSCAN

2. Are clusters roughly spherical?

- Yes → K-Means++ is great
- No → DBSCAN or Hierarchical

3. Are there outliers?

- Yes → DBSCAN (marks as noise)
- No → Any method works

Decision Tree: Which Algorithm to Use?

Choosing a Clustering Algorithm

1. Do you know K ?

- Yes → Consider K-Means++
- No → Consider Hierarchical or DBSCAN

2. Are clusters roughly spherical?

- Yes → K-Means++ is great
- No → DBSCAN or Hierarchical

3. Are there outliers?

- Yes → DBSCAN (marks as noise)
- No → Any method works

4. How large is your dataset?

- Large ($n > 10000$) → K-Means++ (fastest)
- Medium ($1000 \leq n \leq 10000$) → DBSCAN or K-Means++

Decision Tree: Which Algorithm to Use?

Choosing a Clustering Algorithm

1. Do you know K ?

- Yes → Consider K-Means++
- No → Consider Hierarchical or DBSCAN

2. Are clusters roughly spherical?

- Yes → K-Means++ is great
- No → DBSCAN or Hierarchical

3. Are there outliers?

- Yes → DBSCAN (marks as noise)
- No → Any method works

4. How large is your dataset?

- Large ($n > 10000$) → K-Means++ (fastest)
- Medium ($1000 \leq n \leq 10000$) → DBSCAN or K-Means++

Practical Recommendations

Default Choice: K-Means++

- Fast, simple, works well in practice
- Use elbow method to find K
- Run multiple times (non-deterministic)

Practical Recommendations

Default Choice: K-Means++

- Fast, simple, works well in practice
- Use elbow method to find K
- Run multiple times (non-deterministic)

When to use alternatives:

- **Arbitrary shapes + outliers** → DBSCAN
- **Need full hierarchy** → Hierarchical
- **Probabilistic model** → Gaussian Mixture Models
- **Varying densities** → HDBSCAN
- **Very large scale** → Mini-Batch K-Means

Practical Recommendations

Default Choice: K-Means++

- Fast, simple, works well in practice
- Use elbow method to find K
- Run multiple times (non-deterministic)

When to use alternatives:

- **Arbitrary shapes + outliers** → DBSCAN
- **Need full hierarchy** → Hierarchical
- **Probabilistic model** → Gaussian Mixture Models
- **Varying densities** → HDBSCAN
- **Very large scale** → Mini-Batch K-Means

General Advice:

Evaluating Clustering Quality

Problem: Unlike supervised learning, we don't have "true" labels!

Evaluating Clustering Quality

Problem: Unlike supervised learning, we don't have "true" labels!

Internal Metrics (no ground truth needed):

- **Silhouette Score:** Measures compactness vs separation

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$

Range: $[-1, 1]$, higher is better

- **Davies-Bouldin Index:** Ratio of within-cluster to between-cluster distances Lower is better
- **Calinski-Harabasz Index:** Ratio of between-cluster to within-cluster variance Higher is better

Evaluating Clustering Quality

Problem: Unlike supervised learning, we don't have "true" labels!

Internal Metrics (no ground truth needed):

- **Silhouette Score:** Measures compactness vs separation

$$s(i) = \frac{b(i) - a(i)}{\max\{a(i), b(i)\}}$$

Range: $[-1, 1]$, higher is better

- **Davies-Bouldin Index:** Ratio of within-cluster to between-cluster distances Lower is better
- **Calinski-Harabasz Index:** Ratio of between-cluster to within-cluster variance Higher is better

External Metrics (if ground truth available):

- **Adjusted Rand Index (ARI):** Similarity to true labeling

Practical Considerations

Feature Scaling for Clustering

Problem: Features with different scales dominate distance calculations!

Feature Scaling for Clustering

Problem: Features with different scales dominate distance calculations!

Example:

- Age: 20-80 years
- Income: \$20,000 - \$200,000
- Distance dominated by income!

Feature Scaling for Clustering

Problem: Features with different scales dominate distance calculations!

Example:

- Age: 20-80 years
- Income: \$20,000 - \$200,000
- Distance dominated by income!

Solution: Always normalize features before clustering!

StandardScaler (Z-score normalization):

$$x_{\text{scaled}} = \frac{x - \mu}{\sigma}$$

Transforms to mean = 0, std = 1

Feature Scaling for Clustering

Problem: Features with different scales dominate distance calculations!

Example:

- Age: 20-80 years
- Income: \$20,000 - \$200,000
- Distance dominated by income!

Solution: Always normalize features before clustering!

StandardScaler (Z-score normalization):

$$x_{\text{scaled}} = \frac{x - \mu}{\sigma}$$

Transforms to mean = 0, std = 1

MinMaxScaler:

$$x_{\text{scaled}} = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

Transforms to range [0, 1]

Feature Scaling for Clustering

Problem: Features with different scales dominate distance calculations!

Example:

- Age: 20-80 years
- Income: \$20,000 - \$200,000
- Distance dominated by income!

Solution: Always normalize features before clustering!

StandardScaler (Z-score normalization):

$$x_{\text{scaled}} = \frac{x - \mu}{\sigma}$$

Transforms to mean = 0, std = 1

MinMaxScaler:

$$x_{\text{scaled}} = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

Transforms to range [0, 1]

High-Dimensional Data

Curse of Dimensionality:

- In high dimensions, all distances become similar!
- Nearest and farthest neighbors have similar distances
- Clustering becomes meaningless

High-Dimensional Data

Curse of Dimensionality:

- In high dimensions, all distances become similar!
- Nearest and farthest neighbors have similar distances
- Clustering becomes meaningless

Solutions:

1. Dimensionality Reduction first:

- PCA: Project to lower dimensions
- t-SNE: Nonlinear reduction (for visualization)
- Autoencoders: Neural network-based

2. Feature Selection:

- Remove irrelevant features
- Use domain knowledge

Handling Categorical Data

Problem: K-Means uses Euclidean distance (assumes numerical features)

Handling Categorical Data

Problem: K-Means uses Euclidean distance (assumes numerical features)

Solutions:

1. One-Hot Encoding:

- Color = {Red, Blue, Green} → 3 binary features
- Then cluster normally
- Can use Hamming distance instead of Euclidean

Handling Categorical Data

Problem: K-Means uses Euclidean distance (assumes numerical features)

Solutions:

1. One-Hot Encoding:

- Color = {Red, Blue, Green} → 3 binary features
- Then cluster normally
- Can use Hamming distance instead of Euclidean

2. K-Modes:

- Variant of K-Means for categorical data
- Uses mode instead of mean
- Uses Hamming distance

Handling Categorical Data

Problem: K-Means uses Euclidean distance (assumes numerical features)

Solutions:

1. One-Hot Encoding:

- Color = {Red, Blue, Green} → 3 binary features
- Then cluster normally
- Can use Hamming distance instead of Euclidean

2. K-Modes:

- Variant of K-Means for categorical data
- Uses mode instead of mean
- Uses Hamming distance

3. K-Prototypes:

- Handles mixed numerical + categorical

Handling Categorical Data

Problem: K-Means uses Euclidean distance (assumes numerical features)

Solutions:

1. One-Hot Encoding:

- Color = {Red, Blue, Green} → 3 binary features
- Then cluster normally
- Can use Hamming distance instead of Euclidean

2. K-Modes:

- Variant of K-Means for categorical data
- Uses mode instead of mean
- Uses Hamming distance

3. K-Prototypes:

- Handles mixed numerical + categorical

Summary

What We Learned Today

1. Fundamentals:

- Unsupervised learning finds patterns without labels
- Clustering groups similar objects together
- Need to define similarity/distance measure

What We Learned Today

1. Fundamentals:

- Unsupervised learning finds patterns without labels
- Clustering groups similar objects together
- Need to define similarity/distance measure

2. K-Means:

- Minimizes within-cluster sum of squares
- Two-step: Assignment → Update → Repeat
- K-Means++ initialization helps a lot!
- Use elbow method to choose K

What We Learned Today

1. Fundamentals:

- Unsupervised learning finds patterns without labels
- Clustering groups similar objects together
- Need to define similarity/distance measure

2. K-Means:

- Minimizes within-cluster sum of squares
- Two-step: Assignment → Update → Repeat
- K-Means++ initialization helps a lot!
- Use elbow method to choose K

3. Hierarchical:

- Builds tree of clusters (dendrogram)

What We Learned Today (cont.)

4. DBSCAN:

- Density-based: finds arbitrarily shaped clusters
- Identifies outliers as noise
- Parameters: ε (radius), MinPts (threshold)
- Use k-distance graph to choose ε

What We Learned Today (cont.)

4. DBSCAN:

- Density-based: finds arbitrarily shaped clusters
- Identifies outliers as noise
- Parameters: ε (radius), MinPts (threshold)
- Use k-distance graph to choose ε

5. Practical Tips:

- Always scale/normalize features first!
- Visualize results (PCA for high-D)
- Try K-Means++ as baseline
- Choose algorithm based on data properties
- Use multiple metrics to evaluate

What We Learned Today (cont.)

4. DBSCAN:

- Density-based: finds arbitrarily shaped clusters
- Identifies outliers as noise
- Parameters: ε (radius), MinPts (threshold)
- Use k-distance graph to choose ε

5. Practical Tips:

- Always scale/normalize features first!
- Visualize results (PCA for high-D)
- Try K-Means++ as baseline
- Choose algorithm based on data properties
- Use multiple metrics to evaluate

Next Steps

Practice:

- Implement K-Means from scratch (HW/Lab)
- Try scikit-learn: KMeans, DBSCAN, AgglomerativeClustering
- Experiment with real datasets (Iris, customer data, images)

Next Steps

Practice:

- Implement K-Means from scratch (HW/Lab)
- Try scikit-learn: KMeans, DBSCAN, AgglomerativeClustering
- Experiment with real datasets (Iris, customer data, images)

Further Topics:

- Gaussian Mixture Models (probabilistic clustering)
- Spectral clustering (graph-based)
- Topic modeling (Latent Dirichlet Allocation)
- Deep clustering (autoencoders + K-Means)

Next Steps

Practice:

- Implement K-Means from scratch (HW/Lab)
- Try scikit-learn: KMeans, DBSCAN, AgglomerativeClustering
- Experiment with real datasets (Iris, customer data, images)

Further Topics:

- Gaussian Mixture Models (probabilistic clustering)
- Spectral clustering (graph-based)
- Topic modeling (Latent Dirichlet Allocation)
- Deep clustering (autoencoders + K-Means)

Code and Resources:

[Google Colab Notebook](#)

Questions?